

# Introduction to Modern AI

## Week 2: Supervised Learning I - Differentiable Parametric Models

Gavin Hartnett

PRGS, Winter Quarter 2022

# Topics we will cover this week

- In this week we explore more complex models of the form  $f(\mathbf{x}; \boldsymbol{\theta})$
- Mainly in supervised learning context, but many results will apply to other learning paradigms
- As we move through the course  $f(\mathbf{x}; \boldsymbol{\theta})$  will become more complicated (expressive), but it will retain two important properties
  - it will be described in terms of parameters  $\boldsymbol{\theta}$
  - it will be differentiable, i.e.  $\nabla_{\boldsymbol{\theta}} f(\mathbf{x}; \boldsymbol{\theta})$  will exist and be computable
- Main goal of this week is to introduce neural networks and gradient descent
- First, we will warm-up with the Perceptron
- Next week we will look at non-parametric models such as k-nearest neighbors and decision trees

# The Perceptron

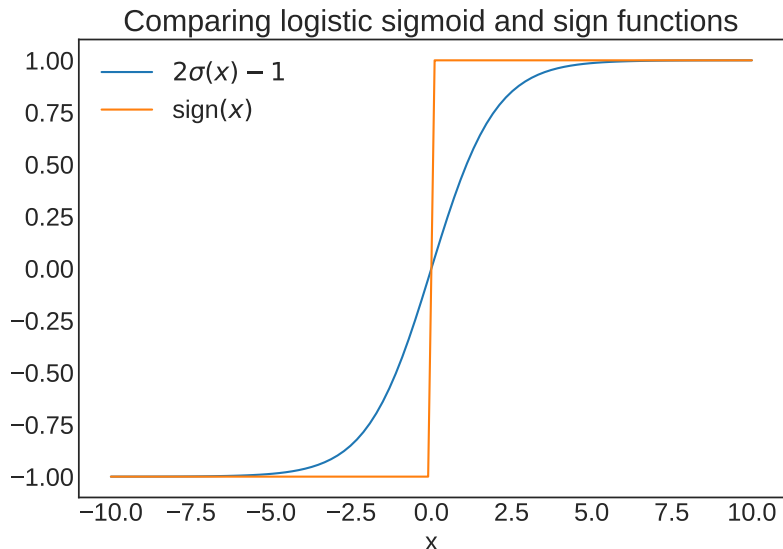
# The Perceptron

- Historically significant model + algorithm for binary classification
- Useful to formulate it using the convention that  $y \in \{-1, +1\}$
- Precursor to the sophisticated neural networks we use today
- Similar to logistic regression
- Linear model (absorbing bias into  $\mathbf{w}$ ):

$$f(\mathbf{x}; \mathbf{w}) = \text{sign}(\mathbf{w}^T \mathbf{x}) = \begin{cases} -1, & \mathbf{w}^T \mathbf{x} < 0 \\ 1, & \mathbf{w}^T \mathbf{x} > 0 \end{cases}$$

- Unlike logistic regression, the perceptron just predicts a class, not a probability

# The Perceptron



# The Perceptron

---

## Algorithm 1 Perceptron Learning Algorithm

---

```
1: initialize weights  $\mathbf{w}$ 
2: while not converged do
3:   for  $i = 1, \dots, N$  do
4:     compute model prediction  $\hat{y}_i = f(\mathbf{x}_i; \mathbf{w})$ 
5:      $\mathbf{w} \leftarrow \frac{\lambda}{2}(\hat{y}_i - y_i)\mathbf{x}_i$ 
6:   end for
7: end while
```

---

Interpretation:

- if prediction is correct, don't change weight vector
- if prediction is wrong, increase/decrease weight vector proportional to  $\mathbf{x}_i$

# The Perceptron

Show Example


# The Perceptron

- The Perceptron learning algorithm is trying to find a separating hyper-plane
- If the data is not linearly separable, then the algorithm will not converge
- There are many solutions when the data is separable, and the algorithm doesn't necessarily pick the best one



# The Perceptron

- The Perceptron was actually built as a piece of specialized hardware (Mark I Perceptron)
- That's cool... but why are you telling us this?
  - Non-linear generalizations of Perceptrons are the simplest Artificial Neural Networks (ANNs)
  - The learning algorithm is actually stochastic gradient descent
  - They are closely related to a powerful set of models called Support Vector Machines (SVMs)



# (Artificial) Neural Networks

# The Need For More Expressive Models

- Both linear regression and logistic regression are *linear* (or log-linear) models
- This property severely constrains the types of functions they can fit
- It would be desirable to have models capable of learning arbitrarily complicated functions
  - Requires flexible families of models (i.e., NNs)
  - Requires general learning algorithm/framework (i.e., gradient descent)

# Basis Expansions

- Recall linear regression model:

$$f(\mathbf{x}; \mathbf{w}, b) = \mathbf{w}^T \mathbf{x} + b$$

- How can this be improved to incorporate non-linearities?
  - Suppose we knew that the probability explicitly depended on  $(x_4)^2$ , in addition to a linear dependence on  $\mathbf{x}$ ?
  - Just add it in!

$$\mathbf{x}' = (x_1, x_2, \dots, x_p, (x_4)^2)$$

$$\mathbf{w}' = (w_1, w_2, \dots, w_p, w_{p+1})$$

- Model remains linear in enlarged  $\mathbf{x}'$  space!
- Suppose it also depends on  $\sin(x_2)$ ?
- Just add it in!

$$\mathbf{x}' = (x_1, x_2, \dots, x_p, (x_4)^2, \sin(x_2))$$

$$\mathbf{w}' = (w_1, w_2, \dots, w_p, w_{p+1}, w_{p+2})$$

# Basis Expansions

- This basic idea can be extended to incorporate arbitrary non-linearities in a systematic fashion
- Suppose we identify a set of  $M$  useful transformations of the original input (or feature) space  $\mathbf{x}$ :  $h_m(\mathbf{x}) : \mathbb{R}^p \rightarrow \mathbb{R}$ , for  $m = 1, \dots, M$
- Simply build a linear model in the transformed feature space:

$$f(\mathbf{x}; \boldsymbol{\theta}) = \sum_{m=1}^M w_m h_m(\mathbf{x}) + b$$

- Issues with this approach
  - Poor scaling:  $p^d$  independent terms (and parameters) for  $d$ -degree polynomial
  - Inefficient way to make model more flexible unless we have special knowledge of the non-linearities

# Neural Networks

- Basic idea is very simple: build expressive model  $f(\mathbf{x}; \theta)$  using simpler modular components called layers:

$$f = f^L \circ f^{L-1} \circ \dots \circ f^1$$


- here  $\circ$  means “compose”, i.e.  $(f \circ g)(x) = f(g(x))$
- the output of layer  $f^{(\ell)}$  becomes the input to layer  $f^{(\ell+1)}$
- As we will see,  $f$  has the structure of a bunch of simple processing units connected through a network - loose inspiration from (real) neural networks
- This lecture: simple, “vanilla” neural network, aka multi-layer perceptron (MLP)
- In later lectures we will learn about many sophisticated variants and extensions

# Neural Networks

- Neural networks are mathematical models,  $f(\mathbf{x}; \theta)$ , where

$$f = f^L \circ f^{L-1} \circ \dots \circ f^1$$

- Can also be described graphically
- Circles represent both variables and a computation, such as
$$y = g(w_1 x_1 + w_2 x_2 + w_3 x_3 + b)$$
- Computation involves add/multiply and a possible non-linearity, called the activation function  $g$
- Lines represent flow of information
- Network is organized into layers, input, hidden ( $\times 2$ ), and output



# Neural Networks

Here is how to “translate” the diagram

- Input (yellow dots):

$$\mathbf{x}^{(0)} = (x_1^{(0)}, x_2^{(0)}, x_3^{(0)})$$

- Hidden Layer 1 (blue dots):


$$\mathbf{x}^{(1)} = (x_1^{(1)}, x_2^{(1)}, x_3^{(1)}, x_4^{(1)}) = g^{(1)} \left( \mathbf{w}^{(1)} \mathbf{x}^{(0)} + \mathbf{b}^{(1)} \right)$$

- Hidden Layer 2 (green dots):

$$\mathbf{x}^{(2)} = (x_1^{(2)}, x_2^{(2)}, x_3^{(2)}, x_4^{(2)}) = g^{(2)} \left( \mathbf{w}^{(2)} \mathbf{x}^{(1)} + \mathbf{b}^{(2)} \right)$$

- Output Layer (red dot):

$$x^{(3)} = (x_1^{(3)}) = g^{(3)} \left( \mathbf{w}^{(3)} \mathbf{x}^{(2)} + \mathbf{b}^{(3)} \right)$$






# Neural Networks

Putting it all together:  $y = x_1^{(3)} = f(\mathbf{x}; \boldsymbol{\theta})$ , with

$$x_1^{(3)} = g^{(3)} \left( \mathbf{w}^{(3)} \left( g^{(2)} \left( \mathbf{w}^{(2)} \left( g^{(1)} \left( \mathbf{w}^{(1)} \mathbf{x}^{(0)} + \mathbf{b}^{(1)} \right) \right) + \mathbf{b}^{(2)} \right) \right) + \mathbf{b}^{(3)} \right)$$

Some comments

- The input and output dimensions are fixed by the problem
- Hidden layer dimensions are unconstrained
- Number of hidden layers is unconstrained
- Each layer can have a different activation function




# Neural Networks

Putting it all together:  $y = x_1^{(3)} = f(\mathbf{x}; \boldsymbol{\theta})$ , with

$$x_1^{(3)} = g^{(3)} \left( \mathbf{w}^{(3)} \left( g^{(2)} \left( \mathbf{w}^{(2)} \left( g^{(1)} \left( \mathbf{w}^{(1)} \mathbf{x}^{(0)} + \mathbf{b}^{(1)} \right) \right) + \mathbf{b}^{(2)} \right) \right) + \mathbf{b}^{(3)} \right)$$


How many parameters does the network contain?

- bias *vectors*
  - $\dim(\mathbf{b}_1) : N_{h_1}$
  - $\dim(\mathbf{b}_2) : N_{h_2}$
  - $\dim(\mathbf{b}_3) : 1$
- weight *matrices*
  - $\dim(\mathbf{w}^{(1)}) : N_{h_1} \times p$
  - $\dim(\mathbf{w}^{(2)}) : N_{h_2} \times N_{h_1}$
  - $\dim(\mathbf{w}^{(3)}) : 1 \times N_{h_2}$
- The model can be made arbitrarily large by increasing the number of units in either hidden layer



# Activation Functions

- Non-linearities come entirely from the activation functions
- Some common choices for  $g$ :
  - Logistic sigmoid:  $g(x) = \sigma(x)$
  - ReLU:  $g(x) = \max(0, x)$
  - tanh:  $g(x) = \tanh(x)$
  - GELU:  $g(x) = \frac{x}{2} \left( 1 + \operatorname{erf} \left( \frac{x}{\sqrt{2}} \right) \right)$



A cartoon drawing of a biological neuron (left) and its mathematical model (right).

# Universal Approximation Theorem

What class of functions can NNs represent?

- By adding more hidden layers (increasing depth) or adding more units to existing layers (increasing width) we can make the model more expressive (contain more parameters)
- More parameters loosely implies a capacity to approximate a larger class of functions
- The universal approximation theorems formally guarantee that, with enough parameters, NNs are capable of approximating any function with arbitrary accuracy
- There are many theorems for different activation functions and assumptions about model architecture (arbitrary width or depth)
- Theorems are not prescriptive - they do not tell you how many units/layers you need to achieve a certain accuracy
- In the worst-case, you often need an exponentially large network
- How to train a NN is an entirely separate matter


# Gradient Descent

# Gradient Descent

- We've already seen that more complex models can be harder to train (fit) than simpler ones
  - The best-fit parameters for linear regression can be solved for in closed-form
  - Fitting logistic regression requires numerical optimization
  - No surprise: NNs will also require numerical optimization
- There are many optimization algorithms, which one(s) should we use?
  - Want a flexible method
  - No assumptions on geometry of loss function (i.e., cannot assume convexity)
  - Want it to be relatively fast/efficient
- Gradient Descent (GD) and variations are the dominant method used for training NNs

# Gradient Descent

- GD is a very simple algorithm that only requires local information:
  - Starting at a given point, find the direction of steepest descent (given by the gradient vector)
  - Then take a small step in this direction
  - Repeat
- Think of a blind man trying to get to the bottom of a hill



# Gradient Descent Algorithm

---

## Algorithm 2 Gradient Descent

---

- 1: function to be minimized:  $f(\mathbf{x})$
  - 2: set learning rate  $\alpha$
  - 3: initialize variable  $\mathbf{x}$
  - 4: **while** not converged **do**
  - 5:    $\mathbf{x} \leftarrow \mathbf{x} - \alpha \nabla f(\mathbf{x})$
  - 6: **end while**
-



# Gradient Descent Algorithm

- Example: Rosenbrock (banana) function

$$f(x, y) = (1 - x)^2 + 10(y - x^2)^2,$$

- The gradient is


$$\nabla f = \begin{bmatrix} -2(1 - x) - 40x(y - x^2) \\ 20(y - x^2) \end{bmatrix}.$$

- The gradient vanishes for the single point  $(x, y) = (1, 1)$ , which is also the global minimum. At this point the function takes on the value  $f(1, 1) = 0$ .

# Gradient Descent Algorithm

- Example: Rosenbrock function


$$f(x, y) = (1 - x)^2 + 10(y - x^2)^2,$$



# Gradient Descent Algorithm

- Example: Rosenbrock function


$$f(x, y) = (1 - x)^2 + 10(y - x^2)^2,$$



# Gradient Descent Algorithm

- Example: Rosenbrock function

$$f(x, y) = (1 - x)^2 + 10(y - x^2)^2,$$



# Gradient Descent Algorithm

Some issues to keep in mind

- The algorithm only converges if the learning rate  $\alpha$  is small enough
- “Small enough” is problem dependent
- GD can get stuck in local minima
- Not well-suited for problems with widely-varying length scales (i.e., some very steep directions and some very shallow directions)

## Example: Quadratic Loss

- To better understand GD, let's examine a simple optimization problem that we can solve exactly
- Consider a quadratic loss function:

$$L(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{A} \mathbf{w} - \mathbf{b}^T \mathbf{w}$$

- $\mathbf{w}, \mathbf{b} \in \mathbb{R}^d$ , and  $\mathbf{A} \in \mathbb{R}^{d \times d}$ . also assume  $\mathbf{A}$  is symmetric and invertible, meaning  $\mathbf{A}^{-1}$  exists
- Gradient is:

$$\nabla_{\mathbf{w}} L = \mathbf{A} \mathbf{w} - \mathbf{b}$$

- Optimal solution is then  $\mathbf{w}^* = \mathbf{A}^{-1} \mathbf{b}$

## Example: Quadratic Loss

- Quadratic loss function:  $L(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{A} \mathbf{w} - \mathbf{b}^T \mathbf{w}$
- Gradient:  $\nabla_{\mathbf{w}} L = \mathbf{A} \mathbf{w} - \mathbf{b}$
- Optimal solution:  $\mathbf{w}^* = \mathbf{A}^{-1} \mathbf{b}$
- Diagonalize:

$$\mathbf{A} = \mathbf{O}^{-1} \mathbf{\Lambda} \mathbf{O}, \quad \mathbf{\Lambda} = \text{diag}(\lambda_1, \lambda_2, \dots, \lambda_d).$$

$$\mathbf{b}' = \mathbf{O} \mathbf{b}, \quad \mathbf{x} = \mathbf{O} \mathbf{w}.$$

- Gradient descent update rule:

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \alpha \nabla_{\mathbf{w}} L(\mathbf{w}^t) = \mathbf{w}^t - \alpha (\mathbf{A} \mathbf{w}^t - \mathbf{b})$$

- Becomes:

$$\mathbf{x}^{t+1} = \mathbf{x}^t - \alpha (\mathbf{\Lambda} \mathbf{x}^t - \mathbf{b}')$$

- In terms of components:

$$x_a^{t+1} = (1 - \alpha \lambda_a) x_a^t + \alpha b'_a$$

## Example: Quadratic Loss

- GD update rule in  $\mathbf{x}$  variable:

$$x_a^{t+1} = (1 - \alpha\lambda_a) x_a^t + \alpha b'_a$$

- Can “unroll”:

$$\begin{aligned} x_a^{t+1} &= (1 - \alpha\lambda_a) ((1 - \alpha\lambda_a) x_a^{t-1} + \alpha b'_a) + \alpha b'_a \\ &= (1 - \alpha\lambda_a)^2 x_a^{t-1} + (1 + (1 - \alpha\lambda_a)) \alpha b'_a \end{aligned}$$

- Continuing:

$$x_a^{t+1} = (1 - \alpha\lambda_a)^{t+1} x_a^0 + \alpha b'_a \sum_{k=0}^t (1 - \alpha\lambda_a)^k$$

- For this to converge as  $t \rightarrow \infty$ , require  $|1 - \alpha\lambda_a| < 1$ . Then first term vanishes and second term is a geometric series, can be summed to give  $\mathbf{x}^* = \mathbf{O}\mathbf{w}^* = \mathbf{\Lambda}^{-1}\mathbf{b}$ :

$$x_a^* = \lim_{t \rightarrow \infty} x_a^t = \alpha b'_a \left( \frac{1}{1 - (1 - \alpha\lambda_a)} \right) = \frac{b'_a}{\lambda_a}$$



## Example: Quadratic Loss

To summarize

- Optimization problem for  $d$ -dim vector  $\mathbf{w}$  decomposes into  $d$  separate 1-dim optimization problems if we work in the eigenbasis
- Convergence requires  $|1 - \alpha\lambda_a| < 1$
- Corresponds to  $0 < \alpha\lambda_a < 2$ 
  - requires all eigenvalues to have same sign - positive for a (global) minimum
  - then  $\alpha > 0$
  - $\alpha < 0$ ,  $\lambda_a < 0$  corresponds to gradient ascent
- Convergence in this case is *exponential*
- But convergence rate is different for each eigen-direction, overall rate is limited by worst eigenvalue

# Stochastic Gradient Descent Algorithm

Let's use gradient descent to minimize a loss function

- Loss function is averaged over training set:

$$\text{loss}(\boldsymbol{\theta}) = \frac{1}{N} \sum_{i=1}^N \ell(\mathbf{x}_i, y_i; \boldsymbol{\theta})$$

- $\ell$  could be the sum of squared errors (SSE), or binary cross entropy (BCE), or really any loss
- Gradient of loss will also be averaged:

$$\nabla_{\boldsymbol{\theta}} \text{loss}(\boldsymbol{\theta}) = \frac{1}{N} \sum_{i=1}^N \nabla_{\boldsymbol{\theta}} \ell(\mathbf{x}_i, y_i; \boldsymbol{\theta})$$

# Stochastic Gradient Descent Algorithm

$$\nabla_{\theta} \text{loss}(\theta) = \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} \ell(\mathbf{x}_i, y_i; \theta)$$

- This expression is not always convenient to work with
- GD requires many iterations, and each iteration requires that the gradient be computed  $N$  times
- $N$  is often large (and we want  $N$  to be as large as possible)
- Idea: let's cut a corner and instead of averaging the gradient over all  $N$  data points, let's average it over a “mini-batch” of  $N_B \ll N$  data points
- Each GD update can use a different, randomly chosen batch
- Note on terminology:
  - $N_B = N$ : gradient descent
  - $1 < N_B < N$ : mini-batch stochastic gradient descent
  - $N_B = 1$ : stochastic gradient descent (aka on-line learning)

# Introduction to Backpropagation

- To use SGD, we must be able to compute the gradient of the loss function  $\nabla_{\theta} \ell$
- Let's consider the MSE loss for a simple NN with 1 hidden layer:

$$\ell(\mathbf{x}_i, y_i; \boldsymbol{\theta}) = (y_i - f(\mathbf{x}; \boldsymbol{\theta}))^2 ,$$

$$f(\mathbf{x}; \boldsymbol{\theta}) = g^{(2)} \left( \mathbf{w}^{(2)} \left( g^{(1)} \left( \mathbf{w}^{(1)} \mathbf{x}^{(0)} + \mathbf{b}^{(1)} \right) \right) + \mathbf{b}^{(2)} \right)$$

- Need to compute  $\nabla_{\mathbf{w}^{(1)}} f, \nabla_{\mathbf{w}^{(2)}} f, \nabla_{\mathbf{b}^{(1)}} f, \nabla_{\mathbf{b}^{(2)}} f$
- This is do-able, but a bit of a pain
- What if we later want to tweak the network? What if we want to consider a large model with 10's or 100's of layers?
- It's desirable to have an automated procedure to compute the gradient for us

# Introduction to Backpropagation

- Suppose we have functions  $f, g$ , and the composition  $h = g \circ f$
- The chain rule let's use calculate the derivative  $h'$ :

$$h' = (g \circ f)' = (g' \circ f)f'$$

or,

$$h'(x) = g'(f(x))f'(x)$$

- Example:

$$\frac{d}{dx} \sin(x)^2 = 2 \sin x \cos x$$

- $f(x) = \sin(x), g(x) = x^2$
- $f'(x) = \cos(x), g'(x) = 2x$
- $h'(x) = (2 \sin x) \cos x$

# Introduction to Backpropagation

- What about functions built using more than 1 composition?

$$f = f^{(L)} \circ f^{(L-1)} \circ \dots \circ f^{(1)}$$

- Apply the chain rule iteratively:

$$f = f^{(L)} \circ F^{(L-1)}, \quad F^{(L-1)} = f^{(L-1)} \circ \dots \circ f^{(1)}$$

$$f' = (f^{(L)'} \circ F^{(L-1)})F^{(L-1)'}$$

- and so on for  $F^{(L-1)'}$ , and then  $F^{(L-2)'}$ , ...
- This “algorithm” can be used to compute the gradient of an arbitrarily complicated neural net
- Q: why is it called backpropagation?
- In week 4 we'll discuss how this is implemented in modern deep learning software libraries using automatic differentiation

# Recap

- Vanilla, Feed-Forward Neural Networks
- (Stochastic) Gradient Descent
- Backpropagation