

# Introduction to Modern AI

## Week 3: Supervised Learning II - Non-parametric models, Ensembles, and Validation

Gavin Hartnett

PRGS, Winter Quarter 2022

# Topics we will cover this week

- Non-parametric models
  - linear regression, logistic regression, and NNs are all examples of parametric models
  - There are many other important parametric models in addition to these
  - There is also a large collection of non-parametric models such as k-nearest neighbors and tree-based models
  - We will remain in the supervised learning setting
- Validation
  - We will discuss methods for validating the performance of a trained model.
  - Again, this applies to both parametric and non-parametric models, but we will focus on non-parametric models here.
- Ensembles
  - It is often desirable to aggregate or pool many individual models to form an ensemble of models
  - This can be done for parametric or non-parametric models, here we will primarily focus on ensembles of the latter

# Review HW1 Solution

Go over the solutions in class

# Non-Parametric Models: kNN and Trees

# $k$ -NN

- Can be used for regression or classification
- Regression
  - The prediction at a point  $\mathbf{x}$  is just the average value of the  $k$  closest points to  $\mathbf{x}$  in the dataset

$$\hat{y}(\mathbf{x}) = \frac{1}{k} \sum_{\mathbf{x}_i \in N_k(\mathbf{x})} y_i$$

- These  $k$ -closest points are called the neighborhood  $N_k(\mathbf{x})$
- Classification
  - The estimate for the probability that the class is  $c$  is just the fraction of points in the neighborhood that are class  $c$

$$\hat{p}(y = c|\mathbf{x}) = \frac{1}{k} \sum_{\mathbf{x}_i \in N_k(\mathbf{x})} \delta_{y(\mathbf{x}), c}$$

## $k$ -NN

- $k$  is a hyper-parameter
  - We are free to choose it
  - not directly optimized over, model is not differentiable wrt to it
  - $k$  is unrelated to number of classes  $K$
- The existence of a neighborhood function  $N_k(\mathbf{x})$  presupposes a distance metric
  - Euclidean distance ( $L^2$  norm):

$$\|\mathbf{x} - \mathbf{x}'\|_2 = \sqrt{\sum_{i=1}^d (x_i - x'_i)^2}$$


- Taxicab or Manhattan distance ( $L^1$  norm):

$$\|\mathbf{x} - \mathbf{x}'\|_1 = \sum_{i=1}^d |x_i - x'_i|$$


- More generally,  $L^p$  norm:

$$\|\mathbf{x} - \mathbf{x}'\|_p = \sqrt[p]{\sum_{i=1}^d |x_i - x'_i|^p}$$

# Aside $L^p$ Norms




# $k$ -NN: Classification Example






# $k$ -NN: Classification Example




# $k$ -NN: Classification Example

- Here  $k = 5$
- distance is Euclidean
- Accuracy is 100% (!)
- Compared to logistic regression
  - Decision boundary is non-linear, appears to be very flexible
  - Model has no parameters - predictions determined directly from dataset



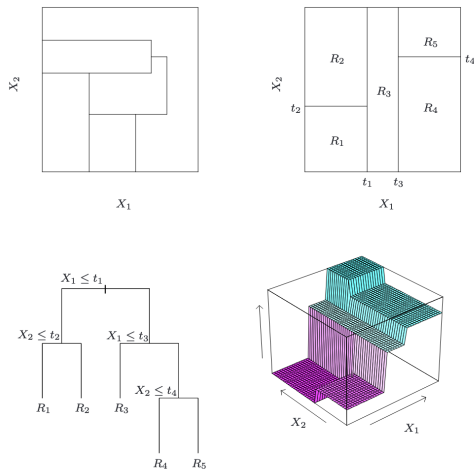
## $k$ -NN: Classification Example #2



# Tree-Based Models

- Tree-based methods are another powerful class of non-parametric models
- Such methods form the basis for many top-performing ML model, such as XGBoost
- We will discuss one popular method for dealing with trees, Classification And Regression Trees (CART), though others exist
- Single-tree models have drawbacks, but can become quite powerful when used in conjunction with ensembling methods (random forests)
- Next few slides closely follow Sec 9.2 of *Elements of Statistical Learning*

# Tree-Based Models



**FIGURE 9.2.** Partitions and CART. Top right panel shows a partition of a two-dimensional feature space by recursive binary splitting, as used in CART, applied to some fake data. Top left panel shows a general partition that cannot be obtained from recursive binary splitting. Bottom left panel shows the tree corresponding to the partition in the top right panel, and a perspective plot of the prediction surface appears in the bottom right panel.

# Tree-Based Models

- Basic idea: use a *binary* tree to hierarchically split input domain into rectangles
- Regression: fit a constant to each domain

$$f(x) = \sum_{m=1}^M c_m I(x \in R_m), \quad \hat{c}_m = \text{ave}(y_i | x_i \in R_m)$$

- Classification: assign a class  $k(m)$  for each region  $m = 1, \dots, M$ . Let  $k(m)$  be the most popular class for training data in region  $R_m$ .
- Main question: how to fit (grow) the tree to data?

# Tree-Based Models

How to grow a regression tree

- Scan over data dimension  $j$  (integer) and split point  $s$  (real number)
- These separate the data into two distinct regions:

$$R_1(j, s) = \{\mathbf{x} | x_j \leq s\}, \quad R_2(j, s) = \{\mathbf{x} | x_j > s\}.$$

- Fit the constants  $\hat{c}_1, \hat{c}_2$  according to  $\hat{c}_m = \text{ave}(y_i | \mathbf{x}_i \in R_m)$
- Compute the MSE for both regions:

$$\sum_{\mathbf{x}_i \in R_1} (y_i - \hat{c}_1)^2 + \sum_{\mathbf{x}_i \in R_2} (y_i - \hat{c}_2)^2$$

- Find the  $(j, s)$  pair that minimizes this
- Repeat process for each region separately

# Tree-Based Models

## Some comments

- When to stop growing the tree? How large should it be?
- General idea is to grow a large tree first, then prune it by collapsing nodes (merging regions)
  - Larger trees are more expressive, but have higher variance
- Are trees really non-parametric?



# Model Validation, Bias-Variance Trade-Off, Regularization

# Train/Test Split

- Then primary goal in ML is to use data to develop a model that performs well on unseen data (called *generalization*)
- We would like to have some idea of how the model will perform on such unseen data
- The main approach to this problem is to split the data into a train and test set
  - Train set: used to train/fit the model
  - Test set: used to evaluate model performance. Proxy for “unseen” data.
- Generally, more data corresponds to a better model, so we want to maximize  $N_{\text{train}}$ .
- However, having a good estimate for model performance is also important, so want  $N_{\text{test}}$  to be large as well.
- No correct way to split the data, often 70/30 or 80/20 is done.


# Train/Validation/Test Split

- Often the data is actually split into 3 datasets
  - Train set: used to train/fit *multiple* models
  - Validation set: used to select the best model
  - Test set: used to evaluate final model performance. Proxy for “unseen” data.
- Often we train multiple models instead of just 1
  - Ensemble methods
  - Cross-validation
  - Regularization
- Once we use the validation set to select the best model, it has been “burned”
- Want test set to give us as unbiased a view of the final model performance as possible

# Bias-Variance Trade-Off

- We've now observed a few interesting points hinting at a fundamental concept in ML dubbed the *bias-variance trade-off*
- The central goal of supervised learning is to develop a model that correctly captures the signal, but not the noise
  - If only the signal is captured, the model will do well when applied to previously unseen data
  - If the noise is captured too much, the model will not generalize and is said to be overfit
- By making the model more expressive (e.g., by adding more parameters or reducing  $k$  in kNN) the model is capable of representing a wider class of functions (less bias)
- On the other hand, with more expressive models there is a danger that the model will fit the noise as well as the signal (high variance)
- Synonyms: capacity, complexity, expressivity

# Bias-Variance Trade-Off



**FIGURE 2.11.** *Test and training error as a function of model complexity.*

# Bias-Variance Trade-Off


Let's be more precise here

- Assume  $y = f(x) + \epsilon$  is the true relationship between  $x, y$
- Here  $\epsilon$  is a random variable with  $\mathbb{E}[\epsilon] = 0$ ,  $\text{Var}[\epsilon] = \sigma^2$ 
  - Recall that  $\text{Var}[X] = \mathbb{E}[X^2 - (\mathbb{E}[X])^2]$
  - And that  $\mathbb{E}$  just means expectation, i.e.  $\mathbb{E}[f(X)] = \int dx p(x)f(x)$
- Also, let the predictions be given by  $\hat{y}(x) = \hat{f}(x; \mathcal{D})$
- Can show (HW exercise!) that the MSE decomposes into 3 terms:

$$\begin{aligned}\text{MSE}(x) &= \mathbb{E}_{\epsilon} (\hat{y}(x) - y)^2 \\ &= \text{Bias}[\hat{f}(x; \mathcal{D})]^2 + \text{Var}[\hat{f}(x; \mathcal{D})] + \sigma^2\end{aligned}$$

- $\text{Bias}[\hat{f}(x)] = (f(x) - \hat{f}(x; \mathcal{D}))$
- Bias term: represents error due to incorrect model assumptions
- Variance term: represents error due to overfitting
- $\sigma^2$  term: irreducible error due to noise

# Bias-Variance Trade-Off



**Figure:** Source: Wikipedia

# Regularization

- Q: Confronted with the fundamental tension of the bias-variance trade-off, how should we make the best determination of model complexity?
- A: Evaluate a range of models with different complexities on a validation set, and take the model with the best performance on this set
- Q: What's the best way to generate a range of similar models with varying complexities?
- One answer is regularization.



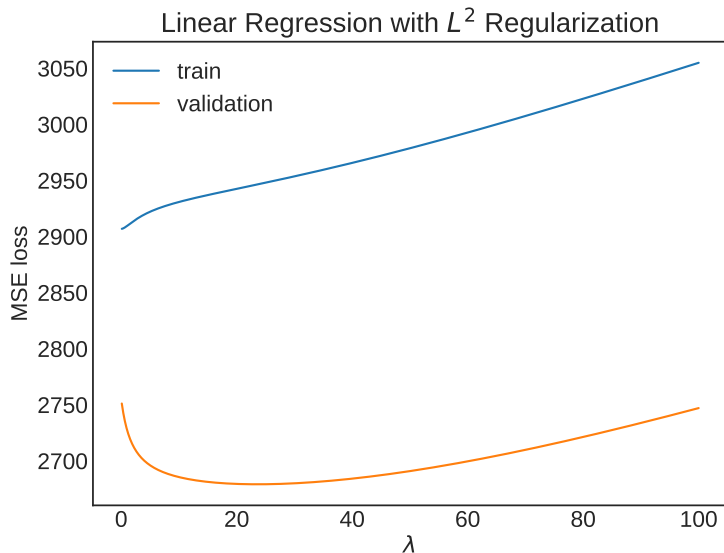
# Regularization via Weight Penalties

- There are many forms of regularization
- For parametric models, a common approach is to use weight penalties

$$\text{loss}'(\mathbf{x}, y; \boldsymbol{\theta}) = \text{loss}(\mathbf{x}, y; \boldsymbol{\theta}) + \lambda \Omega(\boldsymbol{\theta})$$

- $\Omega(\boldsymbol{\theta})$  is zero for  $\boldsymbol{\theta} = 0$ , grows as  $\boldsymbol{\theta}$  grows
- Common examples
  - $\Omega(\boldsymbol{\theta}) = \frac{1}{2} \|\boldsymbol{\theta}\|_2^2$
  - $\Omega(\boldsymbol{\theta}) = \|\boldsymbol{\theta}\|_1$
- $\lambda \geq 0$ : hyper-parameter that controls strength of regularization
- Regularization biases the model to explore less of the parameter space
- Effectively reduces complexity of learned model (complexity shrinks as  $\lambda$  grows)
- Often bias parameters are excluded from regularization

# Regularization Example: Linear Regression



## $L^2$ Regularization

- $L^2$  regularization is also called weight decay, ridge regression
- Loss function:

$$\text{loss}'(\mathbf{x}, y; \boldsymbol{\theta}) = \text{loss}(\mathbf{x}, y; \boldsymbol{\theta}) + \frac{\lambda}{2} \|\boldsymbol{\theta}\|_2^2$$

- Gradient:

$$\nabla_{\boldsymbol{\theta}} \text{loss}'(\mathbf{x}, y; \boldsymbol{\theta}) = \nabla_{\boldsymbol{\theta}} \text{loss}(\mathbf{x}, y; \boldsymbol{\theta}) + \lambda \boldsymbol{\theta}$$

- GD Update:

$$\boldsymbol{\theta}^{(t+1)} = \boldsymbol{\theta}^{(t)} - \alpha \left( \nabla_{\boldsymbol{\theta}} \text{loss}(\mathbf{x}, y; \boldsymbol{\theta}^{(t)}) + \lambda \boldsymbol{\theta}^{(t)} \right) \quad (1)$$

$$= (1 - \alpha\lambda) \boldsymbol{\theta}^{(t)} - \alpha \nabla_{\boldsymbol{\theta}} \text{loss}(\mathbf{x}, y; \boldsymbol{\theta}^{(t)}) \quad (2)$$

- At each step, shrink the weights by a multiplicative factor
- What is the effect over multiple steps?

## $L^2$ Regularization

- Let's revisit quadratic example:

$$L(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{A} \mathbf{w} - \mathbf{b}^T \mathbf{w}$$

- $\mathbf{w}, \mathbf{b} \in \mathbb{R}^d$ , and  $\mathbf{A} \in \mathbb{R}^{d \times d}$ . also assume  $\mathbf{A}$  is symmetric and invertible, meaning  $\mathbf{A}^{-1}$  exists
- Gradient is:

$$\nabla_{\mathbf{w}} L = \mathbf{A} \mathbf{w} - \mathbf{b}$$

- Optimal solution is then  $\mathbf{w}^* = \mathbf{A}^{-1} \mathbf{b}$

## $L^2$ Regularization

- Quadratic loss function:

$$\begin{aligned}L(\mathbf{w}) &= \frac{1}{2} \mathbf{w}^T \mathbf{A} \mathbf{w} - \mathbf{b}^T \mathbf{w} + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w} \\ &= \frac{1}{2} \mathbf{w}^T \mathbf{A}' \mathbf{w} - \mathbf{b}^T \mathbf{w}\end{aligned}$$

- Same as old problem, but with  $\mathbf{A}' = \mathbf{A} + \lambda \mathbb{1}$
- Optimal solution:  $\mathbf{w}^* = (\mathbf{A}')^{-1} \mathbf{b}$
- Diagonalize:

$$\mathbf{A} = \mathbf{O}^{-1} \mathbf{\Lambda} \mathbf{O}, \quad \mathbf{\Lambda} = \text{diag}(\lambda_1, \lambda_2, \dots, \lambda_d).$$

$$\mathbf{b}' = \mathbf{O} \mathbf{b}, \quad \mathbf{x} = \mathbf{O} \mathbf{w}.$$

- Optimal  $x_a^*$  values are shrunk, smaller eigen-directions are shrunk more

$$\mathbf{x}^* = \mathbf{O} \mathbf{w}^* = (\mathbf{\Lambda} + \lambda \mathbb{1})^{-1} \mathbf{b}', \quad x_a^* = \frac{1}{1 + \lambda/\lambda_a} \times \text{old solution}$$

# $L^1$ Regularization

- Called LASSO in statistics
- Loss function:

$$\text{loss}'(\mathbf{x}, y; \boldsymbol{\theta}) = \text{loss}(\mathbf{x}, y; \boldsymbol{\theta}) + \lambda \|\boldsymbol{\theta}\|_1$$

- Gradient:


$$\nabla_{\boldsymbol{\theta}} \text{loss}'(\mathbf{x}, y; \boldsymbol{\theta}) = \nabla_{\boldsymbol{\theta}} \text{loss}(\mathbf{x}, y; \boldsymbol{\theta}) + \lambda \text{sign}(\boldsymbol{\theta})$$

- GD Update:

$$\boldsymbol{\theta}^{(t+1)} = \boldsymbol{\theta}^{(t)} - \alpha \left( \nabla_{\boldsymbol{\theta}} \text{loss}(\mathbf{x}, y; \boldsymbol{\theta}^{(t)}) + \lambda \text{sign}(\boldsymbol{\theta}^{(t)}) \right) \quad (3)$$

- Tends to induce sparsity - small parameters get sent to zero

## $L^1, L^2$ Regularization



**FIGURE 3.11.** Estimation picture for the lasso (left) and ridge regression (right). Shown are contours of the error and constraint functions. The solid blue areas are the constraint regions  $|\beta_1| + |\beta_2| \leq t$  and  $\beta_1^2 + \beta_2^2 \leq t^2$ , respectively, while the red ellipses are the contours of the least squares error function.

# Model Aggregation and Ensembling



# Bootstrapping

- Sometimes data is a scarce resource
- In which case, throwing out 20-25% to do model selection with the validation set, or generalization error estimation with the test set, is quite expensive
- Resampling methods can be used to more efficiently use the limited data you have to estimate various statistical quantities
- Bootstrapping is a popular such method
- Pay attention, because bootstrapping will be important when we consider aggregate/ensemble methods, especially random forests

# Bootstrapping

- Start with a dataset  $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1, \dots, N}$
- By sampling *with replacement*, generate  $B$  new datasets (say,  $B = 100$  or so)
- Note: some data points will be duplicated, some will be absent
- Prob that a given observation will be in a given bootstrap sample:  
 $1 - (1 - 1/N)^N \approx 1 - e^{-1} \approx 0.632$ 
  - single draw:  $P(i) = 1/N$
  - single draw  $P(\neg i) = 1 - 1/N$
  - $P(i \notin \mathcal{D}_b) = (1 - 1/N)^N$
  - $P(i \in \mathcal{D}_b) = 1 - (1 - 1/N)^N$
  - note that

$$e^x = \lim_{N \rightarrow \infty} \left(1 + \frac{x}{N}\right)^N$$

- Train  $B$  distinct models, one for each bootstrap dataset  $\mathcal{D}_b$
- Use these to estimate error or model performance

# Bootstrapping

- Summary:
  - Train  $B$  models on  $B$  independent bootstrap datasets
  - Evaluate model performance by averaging metric of choice across original dataset
  - Can you spot the problem?

# Bootstrapping

- Summary:
  - Train  $B$  models on  $B$  independent bootstrap datasets
  - Evaluate model performance by averaging metric of choice across original dataset
  - Can you spot the problem?
- Each data point used to evaluate the model performance has a 0.632 probability of being in each bootstrap
- Solution: evaluate metric on the so-called out of bag sample
- When evaluating the contribution from datapoint  $i$ , only use bootstrap samples that did not include  $i$

$$\text{bootstrap avg}(h(\hat{y}, y)) = \frac{1}{N} \sum_{i=1}^N \frac{1}{|C^{-i}|} \sum_{b \in C^{-i}} h(f^b(\mathbf{x}_i), y_i)$$

# Bootstrap Aggregating (Bagging)


- Bootstraps were introduced as a validation method
- Goal is not to improve the model, but to improve our understanding of model performance and/or do model selection
- Should be possible to leverage the basic idea to improve our models
- This is achieved by bootstrap aggregating (bagging)
- Regression: average output across bootstrap samples

$$\hat{f}_{\text{Bag}}(\mathbf{x}) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b(\mathbf{x})$$

- Classification: select most popular class

$$\hat{y}_{\text{Bag}}(\mathbf{x}) = \operatorname{argmax}_k \hat{f}_{\text{Bag}}(\mathbf{x}; k)$$

# Bootstrap Aggregating (Bagging)



**Figure:** Source: [https://en.wikipedia.org/wiki/Bootstrap\\_aggregating](https://en.wikipedia.org/wiki/Bootstrap_aggregating)

# Recap

- Introduced general notion of non-parametric models
- As specific examples, introduced kNN and CART
- Introduced validation method, bias-variance trade-off, and weight-norm regularization
- Bootstrapping: take an ensemble of models to reduce variance of the evaluation metric
- Bagging: take an ensemble of models to reduce variance of model predictions
- kNN and trees are generally low-bias, high variance models
- They benefit from model aggregation, but you lose interpretability